

Beyond the Feature dimension

Thus far, the features of an example have been a one-dimensional vector

- the only dimension is the "feature" dimension
- for example, the names of the features in the feature dimension are Price, Volume, Open, Close

But some data has a "shape"

- multiple elements, each having a vector of features
- For example
 - a timeseries: a sequence of elements each having the features Price , Volume, Open, Close
 - a two dimensional image: a grid of elements each have the features Red , Green, Blue


In the diagram

- there are 2 non-feature dimensions, each of length 8
- a (single) feature dimension of length 1

Note: there is **one** feature that appears over a region of size (8×8)

- **not** 64 features arranged in a 2D grid

Two non-feature (spatial) dimensions, single feature



$\mathbf{y}_{(l-1)}$
 $8 \times 8 \times 1$
⏟ ⏟
Spatial Channel

We will start to deal with examples with one or more non-feature dimensions/

- temporal dimension
 - a timeseries with temporal dimension of length 1
 - where $\mathbf{x}_{[t],j}$ denotes feature j at time step t
 - a timeseries with a temporal dimension of length 2
 - $\mathbf{x}_{[d,t],j}$: temporal dimension of length 2 (date d , time t)
- spatial dimension
 - a 2D grid with spatial dimension of length 2
 - where $\mathbf{x}_{[\text{row},\text{col}],j}$ denotes feature j at row row , column col of the grid
- mixed dimension:
 - a timeseries of 2D grids (e.g., a movie) with a temporal dimension of length 1 and spatial dimension of length 2
 - where $\mathbf{x}_{[t,\text{row},\text{col}],j}$ denotes feature j at location (row, col) of frame t of the movie

For clarity

- we have surrounded the non-feature dimensions with brackets
 - writing $\mathbf{x}_{[t],j}$ rather than $\mathbf{x}_{t,j}$
 - rather than a vector as it has been up to this point
- Dropping the brackets
 - \mathbf{x}_t is a vector with n dimensions
 - rather than a scalar (element t of the vector \mathbf{x} whose length is n)

Going forward:

- we may drop the brackets when the context is clear
- our convention is that the *last* dimension in a multi-dimensional object is the *feature* dimension
 - this is called the *channel last* convention
 - when the feature dimension is written as the first dimension: that is called *channel first*
 - TensorFlow layers assume channel last
 - but other toolkits and data providers may not
 - always check !

In this module, we will deal with examples whose features include "spatial" non-feature dimensions.

In a subsequent module, we will deal with examples whose features include "temporal" non-feature dimensions.

The Feature dimension

The feature dimension has some key differences from the non-feature dimensions

- the indices of the feature dimension are *unordered*
 - permuting the features from Price, Volume, Open, Close to Open, Close, Price, Volume does not change the meaning of the example
- the indices of the temporal dimension are totally ordered
 - reversing the indices makes time flow backwards rather than forwards
 - changing the meaning of the example
- the indices of the spatial dimension are (at least, partially) *ordered*
 - there is a spatial relationship between elements whose spatial indices differ by 1
 - row 5 occurs between rows 4 and 6
 - permuting the row order from 4, 5, 6 to 5, 4, 6 changes the meaning of the example

Because of this ordering, the behavior of certain layer types may not respect the order.

For example, consider a sequence of words and its permutation

- $\mathbf{x} = [\text{Machine, Learning, is, easy, not, hard}]$
- $\mathbf{x}' = [\text{Machine, Learning, is, hard, not, easy}]$

The example is a (6×1) vector with a one temporal dimension of length 6 and a single feature word .

- $\mathbf{x}_{[0]}$ is a feature vector of one element: 'Machine'
 - $\mathbf{x}_{[0],0} = \text{'Machine'}$

Clearly, the "meaning" of the two sequences are different.

But suppose we tried to represent this (6×1) vector as vector \mathbf{x}'' of length 6

- no non-feature dimensions
- a single feature dimension of length 6 with features word1, word2, . . . , word6
 - \mathbf{x}'' is a vector of features of length $n = 6$
 - $\mathbf{x}''_0 = \text{'Machine'}$

How would a Fully Connected (Dense) layer deal with the two permutations when they were represented without non-feature dimensions.

It would **not** be able to distinguish between the two permutations !

- A Fully Connected layer computes the dot product of features \mathbf{x} and associated weights \mathbf{w}
 - where \mathbf{x} , \mathbf{w} are vectors of length n containing *only* feature dimensions
- Let perm denote an ordering of indices that is a permutation of $[1, 2, \dots, 6]$
- The dot product of the example and its permutation are the same

$$\mathbf{x} \cdot \mathbf{w} = \mathbf{x}[\text{perm}] \cdot \mathbf{w}[\text{perm}]$$

Our attempt at representing the (6×1) example as 6 features was not successful as it did not respect the ordering of the temporal dimension

Layers transform the *feature* dimension

- transforming raw features into synthetic features
- transforming synthetic features into new synthetic features of increasing complexity

Thus, we need new layer types that

- transform the feature dimension
- for *each element* of the non-feature dimensions

Convolutions

The essence of Neural Networks (as evidenced by the Fully Connected Layer) is pattern matching

- a vector of features from the example
- is matched against a pattern (weights)
- via the dot product
- resulting in a scalar intensity measuring the degree of the match

But the dot product is defined on vectors of features

- no non-feature dimensions
- just a feature dimension of length n

We need to expand the definition of dot product to account for the non-feature dimensions.

We will do this in a subsequent section.

Given a generalized dot product to match features even in the presence of non-feature dimensions

- what would be the best way to use it ?

It would be natural


- to use patterns with non-feature dimensions that were identical in number and length to those of the input.


That is

- we seek to match a pattern against the entire non-feature dimensions of the input.

Consider the following pattern which is of identical dimension to the input

Pattern spanning entire non-feature dimensions, single feature



$\mathbf{y}^{(l-1)}$
 $8 \times 8 \times 1$

Spatial **Channel**

In fact: this pattern is identical to the input and matches it perfectly !

- the generalized dot product of the input and the pattern results in a high activation

But what about inputs similar to this one but

- shifted right/left or up/down
 - we seek "translational invariance"
- a smaller smile
- different distance between the eyes

The pattern would not be as good a match

- lower activation
- even though the "meaning" of the similar input is the same as the original: "smiling face"

It would be more powerful to match patterns against a *sub-region* of the input

- we could look for a small pattern within the input
- while ignoring parts of the input beyond the size of the pattern

For example, suppose

- we want to create an output feature (i.e., match a pattern) that is active when the input contains a pattern matching
 - an "eye"
 - part of the "smile"

We can't use the dot product of the input and a pattern spanning the full input to create these output features

- These patterns are much smaller than the entire input
- The "eye" can occur in any sub-region of the input
- An "eye" can be present regardless of the presence/absence of any other parts of the input


The solution

- define patterns whose non-feature dimensions are *smaller* than that of the entire input
- match these patterns against *every sub-region* of the input of equal size
- producing an output that is a *feature map*
 - an object with non-feature dimensions equal in number and length to that of the input
 - with a single feature
 - that is active at a particular location of the feature map
 - when the pattern matches a sub-region of the input centered at that location

The operation to produce the feature map is called *Convolution*


In the diagram below we show a Convolution that is performed by a CNN Layer type that is layer l of a Sequent NN

- the input is $\mathbf{y}_{(l-1)}$ (the output of layer $(l - 1)$ in a multi-layer NN)
- there are 3 patterns with non-spatial dimensions (2×2)
 - $\mathbf{k}_{(l),1}$ is the pattern for an "eye"
 - $\mathbf{k}_{(l),2}$ and $\mathbf{k}_{(l),3}$ are patterns for the left/right corner of the smile
- the output feature map $\mathbf{y}_{(l)}$ (the layer output)
 - has non-feature dimensions equal in number and length to those of $\mathbf{y}_{(l-1)}$
 - shows the locations within input $\mathbf{y}_{(l-1)}$ where the pattern is matched




$y_{(l-1)}$
 $8 \times 8 \times 1$
 Spatial Channel


Kernels
 (2×2)



$y_{(l),1}$



$y_{(l),2}$



$y_{(l),3}$

$y_{(l)}$
 $8 \times 8 \times 3$
 Spatial Channel

Generalizing the dot product to include non-feature dimensions

We can generalize the dot product to higher dimensions

- let the two arguments \mathbf{x} , \mathbf{k} of the dot product have identical dimensions (both non-feature and feature)
- the point-wise product of the values in corresponding locations of the two arguments
- summed
- to yield a single value

The dot product is

$$\sum_{\mathbf{i} \in I} \mathbf{x}_{\mathbf{i}} * \mathbf{k}_{\mathbf{i}}$$

where I is an enumeration of the index set of \mathbf{x} (and \mathbf{k} since they are of the same dimension)

- each element \mathbf{i} of I is in $D_0 \times \dots \times D_{d-1} \times F$
 - \mathbf{i} is a vector of length $d + 1$
- where there are d non-feature dimensions
 - D_l are the set of indices in non-feature dimension l

This point-wise multiplication is mathematically equivalent to a multiplication described in a way that highlights the feature dimension of the arguments.

- taking sub-vector $\mathbf{x}_{[i]}$ of example \mathbf{x}
 - located a single location \mathbf{i} in the non-feature dimensions of example \mathbf{x} \mathbf{i} is a vector of length N
 - which is a vector of length n features
- and the sub-pattern $\mathbf{k}_{[i]}$ of the pattern \mathbf{k}
 - located at the same location \mathbf{i} in the non-feature dimensions of pattern \mathbf{k}
 - which is a vector of length n features
- taking the dot product

$$\mathbf{x}_{[i]} \cdot \mathbf{k}_{[i]}$$

- adding the dot products computed over *all* the locations i in the non-feature dimensions of \mathbf{x} (and \mathbf{k} since they are of same dimension)

That is, we re-write the point-wise multiplication and sum as

- the dot product at each location in the non-feature dimensions
 - each argument is a vector of n features
- summed

$$\sum_{\mathbf{i} \in D} \mathbf{x}_{[\mathbf{i}]} \cdot \mathbf{k}_{[\mathbf{i}]}$$

where

$$D = D_0 \times \dots \times D_{d-1}$$

is an enumeration of the locations in the non-feature dimensions of \mathbf{x} (and \mathbf{k})

We are matching feature vectors (just as in the one-dimensional case)

- but there is one pair of feature vectors to match for each location in the non-feature dimensions
- so we need to average the pair-wise dot products over the multiple locations

With the addition of non-feature dimensions, we will generalize the pattern matching.

The patterns will now *also* have non-feature dimensions

- same number of dimensions as the example \mathbf{x} 's non-feature dimensions
- but the length f of each dimension in the pattern \mathbf{k} will be *smaller* than the length of the corresponding dimension in the example

Convolution: definition

A convolution involves an example \mathbf{x} and pattern \mathbf{k}

- The *number* of dimensions of \mathbf{x} and \mathbf{k} are the same
- The *length* of the feature dimension of \mathbf{x} and \mathbf{k} are the same
- The *length* of each *non-feature* dimensions of \mathbf{k} is *less than or equal* to the length of the c corresponding non-feature dimension of \mathbf{x}
 - customarily, each non-feature dimension of \mathbf{k} is the same length: f
 - but this is not necessary

The pattern \mathbf{k} , since the length of the non-feature dimensions may be smaller than the corresponding non-feature dimensions of \mathbf{x}

- can only be matched against a sub-region of \mathbf{x} whose non-feature dimensions have length f
 - so that the dimensions of the sub-region match those of the pattern

For example suppose the non-feature dimensions \mathbf{x} and \mathbf{k} are (5×5) and (3×3) respectively.

One can take the dot product of \mathbf{k} with any sub-region of \mathbf{x} of size (3×3) .

There are many such sub-regions of \mathbf{x} .

Convolution is an operation that computes the dot product of \mathbf{k} and each sub-region of \mathbf{x} .

Let there be N non-feature dimensions in \mathbf{x} and \mathbf{k}

- D_l are the set of indices in non-feature dimension l of \mathbf{x}
- D'_l are the set of indices in non-feature dimension l of \mathbf{k}
 - length of each D'_l is f
- F is the set of indices (size n) of the feature dimension of both \mathbf{x} and \mathbf{k}

The index set of

- \mathbf{x} is $(D_0 \times \dots \times D_{d-1} \times F)$
- the non-feature dimensions of \mathbf{x} is

$$D = D_0 \times \dots \times D_{d-1}$$
- \mathbf{k} is $(D'_0 \times \dots \times D'_{d-1} \times F)$
 - length of each D'_l is f
 - $f \leq |D_l|$ for each $0 \leq l \leq (N - 1)$

We will create sub-regions of \mathbf{x} .

Let $\mathbf{x}_{[i]}$ be the sub-region of \mathbf{x} centered at index $i \in D$

- whose non-feature dimensions are of length f
- with feature dimension of length n

The dimensions of $\mathbf{x}_{[i]}$ and \mathbf{k} are identical

- so we can compute the dot product

Let

- $\text{Conv}(\mathbf{x}, \mathbf{k})$ denote the operation performing the convolution of example \mathbf{x} and pattern \mathbf{k}
- let $\mathbf{y} = \text{Conv}(\mathbf{x}, \mathbf{k})$ be the output of the convolution operation
 - the non-feature dimensions of \mathbf{y} and \mathbf{x} are identical in number and length
 - the feature dimension of \mathbf{y} is length 1

Convolution is defined as

$$\mathbf{y}_{[i],0} = \mathbf{x}_{[i]} \cdot \mathbf{k}$$

That is

- the non-feature dimensions of \mathbf{y} and \mathbf{x} are equal in number and length
- location \mathbf{i} in the non-feature dimension of \mathbf{y} is the result of matching
 - the pattern \mathbf{k}
 - against the sub-region of \mathbf{x}
 - centered at \mathbf{i}
 - whose length of non-feature dimensions matches those of \mathbf{k}

Just as a Fully Connected layer can create more than one output feature

- using a different pattern for each output feature
- we can define Convolution to output more than one feature
 - using a different pattern for each feature
 - the patterns are higher-dimensional for the Convolution compare to Fully Connected
 - since the Convolutional patterns contain non-feature dimensions
 - as well as the feature dimension

Let \mathbf{k}_j denote the pattern used for output feature j .

Then

$$\mathbf{y}_{[i],j} = \mathbf{x}_{[i]} \cdot \mathbf{k}_j$$

The output for feature j , denoted as $\mathbf{y}_{\dots,j}$ is called *feature map j*

- shows how each location in the non-feature dimensions of \mathbf{x}
- is matched against pattern \mathbf{k}_j

Examples

It may be easier to grasp the working of a CNN layer

- by starting with simple cases of \mathbf{x} and \mathbf{k}
- and working up to more complex cases
- using pictures to illustrate

We do this in the next module

Where do the patterns come from ? Training a CNN

Hopefully you understand how patterns (kernels) are "feature recognizers".

But you may be wondering: how do we determine the weights in each kernel ?

Answer: a Convolutional Layer is "just another" layer in a multi-layer network

- The kernels are just weights (like the weights in Fully Connected layers)
- We solve for all the weights \mathbf{W} in the multi-layer network in the same way

The answer is: exactly as we did in Classical Machine Learning

- Define a loss function that is parameterized by \mathbf{W} :

$$\mathcal{L} = L(\hat{\mathbf{y}}, \mathbf{y}; \mathbf{W})$$

- The kernel weights are just part of \mathbf{W}
- Our goal is to find \mathbf{W}^* the "best" set of weights

$$\mathbf{W}^* = \underset{\mathbf{W}}{\operatorname{argmin}} L(\hat{\mathbf{y}}, \mathbf{y}; \mathbf{W})$$

- Using Gradient Descent !

In other words: there is nothing special about finding the "best" kernels.

```
In [ ]: print("Done")
```

PREVIOUS VERSION OF THIS NOTEBOOK

The remainder of this notebook is a previous version


- the newer version (sections above) are hopefully improved
- the sections below are kept only for historic reference and will soon be deleted

Convolutional Neural Networks

A Fully Connected/Dense Layer with a single unit producing a single feature at layer l computes

$$\mathbf{y}_{(l),1} = a_{(l)}(\mathbf{y}_{(l-1)} \cdot \mathbf{W}_{(l),1})$$

Fully connected, single feature



That is:

- It recognizes one new synthetic feature
- In the entirety ("fully" connected) of $\mathbf{y}_{(l-1)}$
- Using pattern $\mathbf{W}_{(l),1}$ (same size as $\mathbf{y}_{(l-1)}$)
- To reduce $\mathbf{y}_{(l-1)}$ to a single feature.

The pattern being matched spans the entirety of the input

- Might it be useful to recognize a smaller feature that spanned only *part* of the input ?
- What if this smaller feature could occur *anywhere* in the input rather than at a fixed location ?

For example

- A "spike" in a time series
- The eye in a face

A pattern whose length was that of the entire input could recognize the smaller feature only in a *specific* place

This motivates some of the key ideas behind a Convolutional Layer.

- Recognize smaller features within the whole
- Using small patterns
- That are "slid" over the entire input
- Localizing the specific part of the input containing the smaller feature

The spatial dimension

A small pattern (less than full length of input) can match a sub-section of input at any location.


We can imagine centering the pattern on each input element and applying the match.

The output, therefore, will be a vector of length equal to the input length.

Here is the connectivity diagram of a Convolutional Layer producing a **single** feature at layer l

- Using a pattern of length 3
- Eventually we will show how to produce *multiple* features
- Hence the subscript "1" in $y_{(l),1}$ to denote the first output feature

Convolutional layer, single feature



The output vector $\mathbf{y}_{(l),1}$ is called the first **feature map** as it attempts to match the first feature at each input location.

We refer to the indices of the feature map as the **spatial dimensions**

Thus, the output of layer l is 2 (or higher !) dimensional when layer l is a CNN

- a number of features
 - we have only shown a single feature thus far
- each feature producing a feature map
 - a feature map dimensions are called the spatial dimensions

The output of a CNN layer is a collection of

- $n_{(l)}$ feature maps (one per feature)
- each feature map having the same spatial dimension as its inputs

We can connect multiple CNN layers in sequence

- preserving the spatial dimensions across layers
- but creating more complex features as we get deeper

Technical note: special case

- The input to the first CNN layer l is often a one-dimensional vector of $n_{(l-1)}$ features
- A CNN treats this as a $(n_{(l-1)} \times 1)$ vector
 - 1 feature of a 1D spatial dimensions of shape $((n_{(l-1)},)$

Our convention will be that the **feature dimension** will appear as the **final** dimension of the output of layer l .

- all prior dimensions will be part of the **spatial dimension**

We need to distinguish which dimension is the feature dimension because

- **A convolution finds small patterns in the spatial dimension, not the feature dimension**

To be clear

- the vector of shape $(1 \times d)$ denotes d features at a single spatial location
- the vector of shape $(d \times 1)$ denotes a single features at d spatial locations

Notation

- the feature dimension will be the last index
- $n_{(l)}$ will always denote the *number of features* of a layer l
- $\mathbf{y}_{(l),j',j}$ denotes feature j of layer l at spatial location j'

We say that the above convolutional layer l

- Maps a single feature (defined over a 1D spatial dimension with $d_{(l)} = d_{(l-1)}$ locations) of layer $(l - 1)$
- To a single feature, defined over an identical number of spatial locations in layer l

The importance of the spatial dimension

Let's contrast the CNN layer with a Fully Connected layer.

- The Fully Connected layer we depicted matches a pattern over the full *feature* dimension
 - There is no ordering (or spatial relationship) between features
- The CNN layer we depicted matches a pattern over the full *spatial* dimensions

Spatial dimensions are different than feature dimensions

- They have "order" (spatial relationships)

To see this, we show that a FC layer is insensitive to ordering of inputs

- Consider a vector \mathbf{x} of n features (input to the Fully Connected layer)
- Let perm be permutation of the indices of \mathbf{x} : $[1 \dots n]$.

If we permute both \mathbf{x} and weights Θ , the dot product remains unchanged

$$\Theta^T \cdot \mathbf{x} = \Theta[\text{perm}]^T \cdot \mathbf{x}[\text{perm}]$$

So shuffling inputs to a FC layer does not affect its outputs

- assuming they are shuffled the same way during training and inference

But for certain types of inputs (e.g. images) it is easy to imagine that spatial locality is important.

- Consider a 2D pixel grid depicting a face
- The relative ordering of pixels may be what **defines** a pattern to be recognized
 - The relative location of the pixels within the left eye are important
 - The relative location of the pixels constituting the left eye, right eye, nose and mouth are important

By using a small pattern (and restricting connectivity) we **emphasize the relative locations** of elements

- neighboring elements more important than far away elements.

The "spatial" dimension implies an ordering of elements

- but the ordering does not have to be in space
 - e.g., can be ordered in time

Consider the time series of prices of a single ticker over d days.

Two representations

- $(d \times 1)$: 1 feature ("price") over d spatial ("date") locations
- $(1 \times d)$: 1 ticker with d features (price 1, \dots , price d)

The choice of where the singleton dimension appears is sometimes a matter of interpretation.

- but the last index always denotes the feature dimension

Mathematically, the One Dimensional Convolutional Layer (Conv1d) we have shown computes $\mathbf{y}_{(l)}$

$$\mathbf{y}_{(l),1} = \begin{pmatrix} a_{(l)} \left(N(\mathbf{y}_{(l-1)}, \mathbf{W}_{(l),1}, 1) \cdot \mathbf{W}_{(l),1} \right) \\ a_{(l)} \left(N(\mathbf{y}_{(l-1)}, \mathbf{W}_{(l),1}, 2) \cdot \mathbf{W}_{(l),1} \right) \\ \vdots \\ a_{(l)} \left(N(\mathbf{y}_{(l-1)}, \mathbf{W}_{(l),1}, d_{(l-1)}) \cdot \mathbf{W}_{(l),1} \right) \end{pmatrix}$$

where $N(\mathbf{y}_{(l-1)}, \mathbf{W}_{(l),1}, j)$

- selects a subsequence of $\mathbf{y}_{(l-1), \dots, 1}$ centered at $\mathbf{y}_{(l-1), j, 1}$
 - Note the extra spatial dimension in the subscripting; "... " denotes the full spatial dimension
 - Centered at the j^{th} element in the spatial dimension of feature 1 of layer $(l - 1)$

Note that

- The *same* weight matrix $\mathbf{W}_{(l),1}$ is used for the first feature at *all* locations j
- The size of $\mathbf{W}_{(l),1}$ is the same as the size of the subsequence $N(\mathbf{y}_{(l-1)}, \mathbf{W}_{(l),1}, j)$
 - Since dot product is element-wise multiplication
- The spatial dimension $d_{(l)}$ of $\mathbf{y}_{(l),1}$ is equal to $d_{(l-1)}$

Kernel, Filter

The vector $\mathbf{W}_{(l),1}$ above

- Is a smaller pattern
- That is applied to *each* spatial location j in $\mathbf{y}_{(l-1)}$
- $\mathbf{y}_{(l),j,1}$ recognizes the match/non-match of the smaller first feature at the spatial locations centered at $\mathbf{y}_{(l-1),j,1}$

$\mathbf{W}_{(l),1}$ is called the first convolutional *filter* or *kernel*

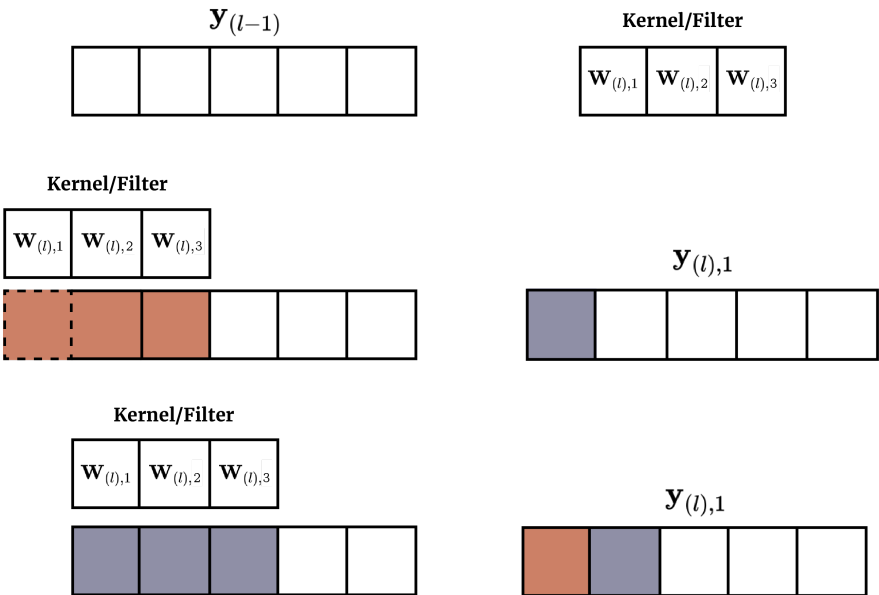
- We will often denote it $\mathbf{k}_{(l),1}$
- But it is just a part of the weights \mathbf{W} of the multi-layer NN.
- We use $f_{(l)}$ to denote the size of the smaller pattern called the *filter size*

A *Convolution* is often depicted as

- A filter/kernel
- That is slid over each location in the input
- Producing a corresponding output for that location

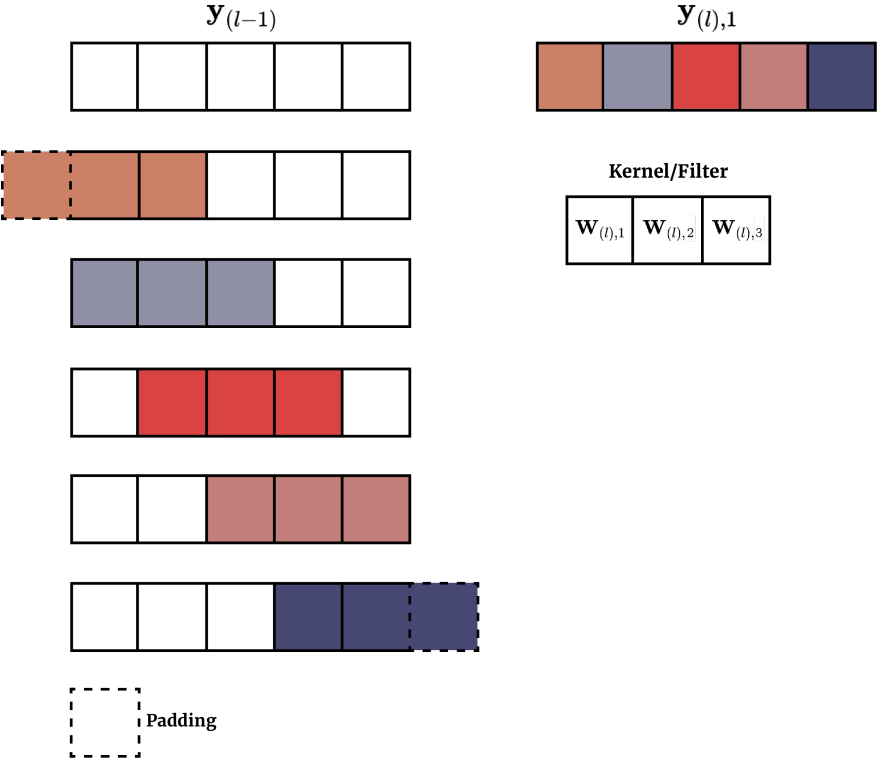
Here's a picture with a kernel of size $f_{(l)} = 3$

Conv 1D, single feature: sliding the filter



After sliding the Kernel over the whole $\mathbf{y}_{(l-1)}$ we get:

Conv 1D, single feature: output feature map



Element j of output $\mathbf{y}_{(l), \dots, 1}$ (i.e., $\mathbf{y}_{(l), j, 1}$)

- Is colored (e.g., $j = 1$ is colored Red)
- Is computed by applying the *same* $\mathbf{W}_{(l), 1}$ to
 - The $f_{(l)}$ elements of $\mathbf{y}_{(l-1), 1}$, centered at $\mathbf{y}_{(l-1), j, 1}$
 - Which have the same color as the output

Note however that, at the "ends" of $\mathbf{y}_{(l-1)}$ the kernel may extend beyond the input vector.

In that case $\mathbf{y}_{(l-1)}$ may be extended with *padding* (elements with 0 value typically)

Activation of a CNN layer

Just like the Fully Connected layer, a CNN layer is usually paired with an activation.

The default activation $a_{(l)}$ in Keras is "linear"

- That is: it returns the dot product input unchanged
- Always know what is the default activation for a layer; better yet: always specify !

Conv2d in action

We have thus far depicted a spatial dimension of length 1.

We can easily expand this into 2 spatial dimensions

- each feature map is 2 dimensional
- each location in the feature map corresponds to a position in a 2D grid

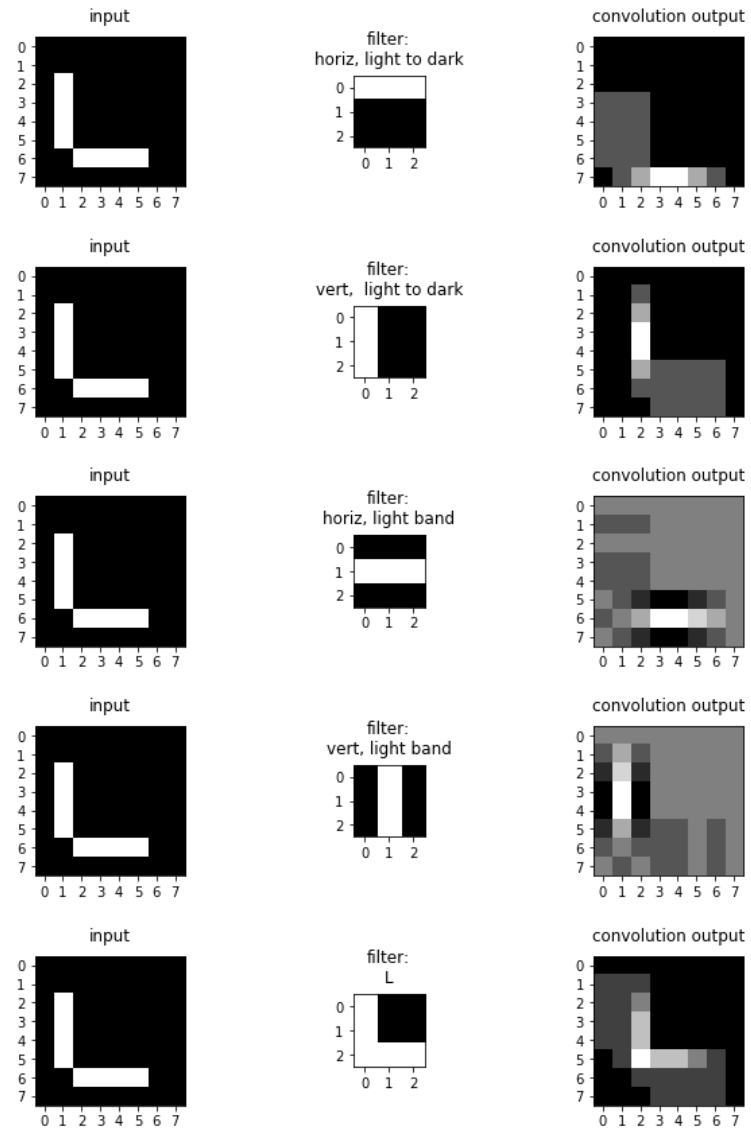
Pre-Deep Learning: manually specified filters have a rich history for image recognition.

Here is a list of manually constructed kernels (templates) that have proven useful

- [list of filter matrices \(https://en.wikipedia.org/wiki/Kernel_\(image_processing\)\)](https://en.wikipedia.org/wiki/Kernel_(image_processing))

Let's see some in action to get a better intuition.

```
In [4]: _ = cnnh.plot_convs()
```



- A bright element in the output indicates a high, positive dot product
- A dark element in the output indicates a low (or highly negative) dot product

In our example

- $N = 2$: Two spatial dimensions
- One input feature: $n_{(l-1)} = 1$
- One output feature $n_{(l)} = 1$
- $f_{(l)} = 3$
 - Kernel is $(3 \times 3 \times 1)$.

The template match will be maximized when

- high values in the input correspond to high values in the matching location of the template
- low values in the input correspond to low values in the matching locations of the template

We can have "spatial" dimensions of length even greater than 2

When we want to emphasize the number of features $n_{(l)}$ rather than the number of spatial dimensions, we will use ellipsis (dots)

$$\mathbf{y}_{\dots, n_{(l)}}$$

where the ellipsis (. . .) is a place-holder for the spatial dimension shape.

In [5]: `print("Done")`

Done

